

optim() in R – Exam Reference Sheet

Contents

| | |
|--|----------|
| 1 Core Usage | 1 |
| 1.1 Function signature | 1 |
| 1.2 Returned object | 2 |
| 2 Objective function pattern | 2 |
| 3 Optimisation methods in optim() | 3 |
| 3.1 "Nelder-Mead" (default) | 3 |
| 3.2 "BFGS" | 3 |
| 3.3 "CG" (Conjugate Gradient) | 3 |
| 3.4 "L-BFGS-B" | 3 |
| 3.5 "SANN" (Simulated Annealing) | 4 |
| 4 Useful control options | 4 |
| 5 Debugging optim() | 4 |
| 5.1 Checklist | 4 |
| 5.2 Typical error messages and fixes | 5 |
| 6 Worked examples | 5 |
| 6.1 Example 1: Linear regression via BFGS | 5 |
| 6.2 Example 2: Same model, with constraint $b \in [0, 2]$ (L-BFGS-B) | 6 |
| 6.3 Example 3: Discrete optimisation with SANN (sketch) | 6 |
| 7 Slide example: “Why does it not work?” | 7 |
| 8 Rapid troubleshooting checklist | 8 |

1 Core Usage

1.1 Function signature

```
optim(par, fn, gr = NULL, ...,
      method = c("Nelder-Mead", "BFGS", "CG", "L-BFGS-B",
                "SANN", "Brent"),
      lower = -Inf, upper = Inf, control = list())
```

Key points:

- **par**: numeric vector of initial parameter values.
- **fn**: **objective function to minimise**. *First argument must be the parameter vector.* Must return a **single numeric value**.
- **gr**: gradient of **fn** (same first argument). Optional. If omitted, finite differences are used for gradient-based methods.

- `...`: extra arguments passed to `fn` and `gr` (typically data).
- `method`: optimisation algorithm (see Section 3).
- `lower`, `upper`: bounds (vectors same length as `par`) for "L-BFGS-B" and "Brent".
- `control`: list of tuning options (Section 4).

1.2 Returned object

```
opt <- optim(...)  
str(opt)
```

Important components:

- `opt$par`: parameter estimate at the found minimum.
- `opt$value`: objective value `fn(opt$par)`.
- `opt$counts`: number of function / gradient evaluations.
- `opt$convergence`: convergence code:
 - 0: successful convergence.
 - 1: iteration limit reached.
 - 10: degenerate simplex (Nelder–Mead).
 - 51, 52: L-BFGS-B warning / error.
- `opt$message`: text message (especially for L-BFGS-B).

Always check `opt$convergence` and `opt$message`.

2 Objective function pattern

General pattern for regression-style problems:

```
crit <- function(par, x, y) {  
  # unpack parameters  
  a <- par[1]  
  b <- par[2]  
  
  # model prediction  
  yhat <- a + b * x  
  
  # objective: sum of squared residuals  
  residuals <- y - yhat  
  sum(residuals^2)  
}
```

Rules:

1. First argument = parameter vector (`par` / `theta`).
2. Return a *single* numeric value (scalar).
3. For maximisation (e.g. log-likelihood), minimise the negative:

```
crit <- function(par, data) {  
  - logLik(par, data)  # minus sign!  
}
```

or use `control = list(fnscale = -1)`.

3 Optimisation methods in optim()

3.1 "Nelder-Mead" (default)

- Derivative-free simplex method.
- Uses only function values, robust to noise / non-smoothness.
- Slower and less reliable in high dimensions, can stagnate (code 10).
- Good first try if gradients are hard to compute.

3.2 "BFGS"

- Quasi-Newton method; builds Hessian approximation.
- Needs gradient (supplied or finite-difference).
- Typically fast and accurate for smooth, unconstrained problems.
- No bounds; use when parameters are effectively unbounded.

Example pattern:

```
optim(par = c(a = 0, b = 0),
      fn  = crit,
      gr  = grad, # optional
      x   = x, y = y,
      method = "BFGS")
```

3.3 "CG" (Conjugate Gradient)

- Uses gradient only, no Hessian stored (low memory).
- Generally slower / more fragile than BFGS, but lighter.
- Unconstrained only. Same calling pattern as BFGS.

3.4 "L-BFGS-B"

- Limited-memory BFGS with box constraints.
- Handles bounds: `lower`, `upper` (vectors, `length = length(par)`).
- Initial `par` must satisfy $\text{lower} \leq \text{par} \leq \text{upper}$.
- Requires *finite* `fn` values inside box; `NA` / `Inf` cause errors.
- Good for parameters like variances, probabilities, rates, etc. with natural bounds.

Example:

```
optim(par    = c(a = 0, b = 1),
      fn     = crit,
      gr     = grad,
      x = x, y = y,
      method = "L-BFGS-B",
      lower  = c(-Inf, 0),    # a unbounded, b >= 0
      upper  = c( Inf, 2))  # b <= 2
```

3.5 "SANN" (Simulated Annealing)

- Stochastic global search (simulated annealing).
- Uses only function values.
- No convergence test; run for fixed number of evaluations.
- Very slow; mainly for rough, multimodal surfaces.
- Often used to get near good region, then refine with BFGS.

Skeleton:

```
optim(par = par0,
      fn = crit,
      method = "SANN",
      control = list(maxit = 10000, temp = 10, trace = 1))
```

4 Useful control options

- **maxit**: maximum iterations (BFGS/CG/L-BFGS-B) or function evaluations (SANN).
- **trace**: if > 0 , prints progress info (debugging).
- **fnscale**:
 - default = 1; if set to -1 , minimising **fn** is equivalent to maximising it.
- **parscale**: scaling factors for parameters; helpful if parameters are on very different scales.
- **REPORT**: how often to print progress when **trace** > 0 .

Example:

```
optim(par = par0, fn = crit, x = x, y = y,
      method = "BFGS",
      control = list(maxit = 1000, trace = 1))
```

5 Debugging optim()

5.1 Checklist

1. **Objective signature**: first argument is parameter vector; returns scalar.
2. **Test fn**: run **fn(par0, ...)** manually. It must be numeric and finite.
3. **Bounds**: when using L-BFGS-B, check

```
all(par0 >= lower & par0 <= upper)
```

4. **Gradient**: if supplying **gr**, confirm dimensions and correctness. If in doubt, remove **gr** and let **optim()** finite-difference it.
5. **Inspect output**: always read **opt\$convergence** and **opt\$message**.

5.2 Typical error messages and fixes

"L-BFGS-B needs finite values of 'fn'"

- Your function returned NA, NaN or Inf.

- Check:

```
fn(par0, ...) # is it finite?
```

- Either:
 - Restrict bounds slightly away from problematic value, or
 - Modify fn to return a large finite penalty instead of Inf.

Convergence code 1: iteration limit reached

- Try increasing maxit or improving starting values.
- Consider a different method (e.g. BFGS instead of Nelder–Mead).

Convergence code 10: degenerate simplex (Nelder–Mead)

- Simplex collapsed; algorithm stagnated.
- Try different initial par, rescaling parameters, or a gradient-based method.

Weird parameter at extreme value

- Often indicates wrong objective function (e.g. summing residuals instead of squared residuals).
- Verify that fn truly matches the quantity you intend to minimise.

6 Worked examples

6.1 Example 1: Linear regression via BFGS

Goal: estimate intercept a and slope b in

$$y_i = a + bx_i + \varepsilon_i$$

by minimising sum of squared residuals.

```
set.seed(1)
x <- 1:20
y <- 1.5 + 2.5 * x + rnorm(length(x), sd = 5)

# objective: sum of squared errors
crit <- function(par, x, y) {
  a <- par[1]; b <- par[2]
  yhat <- a + b * x
  resid <- y - yhat
  sum(resid^2)
}

# analytic gradient (optional but good practice)
grad <- function(par, x, y) {
  a <- par[1]; b <- par[2]
  yhat <- a + b * x
```

```

resid <- y - yhat
d_da <- -2 * sum(resid)      # derivative wrt a
d_db <- -2 * sum(resid * x)  # derivative wrt b
c(d_da, d_db)
}

opt <- optim(par = c(a = 0, b = 0),
             fn = crit,
             gr = grad,
             x = x, y = y,
             method = "BFGS")

opt$par      # estimated (a, b)
opt$value    # SSE at optimum
opt$convergence

```

Key points:

- Objective returns scalar SSE.
- Gradient matches the objective.
- **No bounds**, so we use BFGS.

6.2 Example 2: Same model, with constraint $b \in [0, 2]$ (L-BFGS-B)

```

optB <- optim(par = c(a = 0, b = 0),
               fn = crit,
               gr = grad,
               x = x, y = y,
               method = "L-BFGS-B",
               lower = c(-Inf, 0),
               upper = c( Inf, 2))

optB$par      # (a, b) with b constrained to [0, 2]
optB$convergence

```

If the unconstrained optimum has $b > 2$, the constrained optimum will sit at $b = 2$. Check whether any component of `optB$par` equals its bound to see if a constraint is active.

6.3 Example 3: Discrete optimisation with SANN (sketch)

When the surface is rough / multimodal:

```

crit <- function(par, data) {
  # e.g. negative log-likelihood with many local optima
}

optS <- optim(par = par0,
              fn = crit,
              data = data,
              method = "SANN",
              control = list(maxit = 20000, temp = 10, trace = 1))

# Optionally refine with BFGS
optRef <- optim(par = optS$par,
                 fn = crit,
                 data = data,
                 method = "BFGS")

```

7 Slide example: “Why does it not work?”

From the lecture slide:

```
par(mfrow=c(1,2), font.lab=2, font.axis=2)
model <- function(x, theta){
  return(theta * x)
}

crit <- function(theta, x, y){
  # must have theta as 1st parameter and return a single value...
  return( sum( y - model(x, theta) ) )
}

thbar <- 1.8
x <- rep(c(1,3,7,8), len = 100)
y <- model(x, thbar) + rnorm(x)

plot(x, y, pch=20, cex=1.5, main="optim example 1: data")
points(x, model(x, thbar), col="green", pch=20, cex=2)

optim.out <- optim(par=c(1), fn=crit, x=x, y=y,
                     method="L-BFGS-B", lower=c(0.01), upper=c(3.05))
```

Issue 1: wrong objective (sum of residuals)

The criterion is

$$\text{crit}(\theta) = \sum_i (y_i - \theta x_i).$$

This is *not* a sensible loss:

- Residuals can cancel (positive vs. negative), so the sum can be small even if individual errors are large.
- For these simulated data, the sum is approximately linear in θ , so the minimum lies at one of the *bounds* (0.01 or 3.05), not near the true value 1.8.
- Consequently, `optim()` drives θ to a boundary instead of a meaningful estimate.

Fix

Use **sum of squared residuals**:

```
crit <- function(theta, x, y) {
  resid <- y - model(x, theta)
  sum(resid^2)
}
```

Now the function has a unique minimum near the true $\theta = 1.8$, and methods like "BFGS" or "L-BFGS-B" will find a sensible solution.

Other potential pitfalls in this pattern

- If `crit()` returned a vector, or if the first argument wasn't `theta`, `optim()` would either error or behave unexpectedly.
- If L-BFGS-B is used and the objective becomes `Inf` at a boundary, the call could fail with “L-BFGS-B needs finite values of ‘fn’”.

8 Rapid troubleshooting checklist

When `optim()` “does not work”, check in this order:

1. **Signature:** Is parameter vector first argument in `fn` and `gr`? Does `fn()` return a scalar?
2. **Finite values:** Is `fn(par0, ...)` finite? (Especially for L-BFGS-B.)
3. **Right objective:** Are you minimising the correct quantity (e.g. SSE, negative log-likelihood)?
4. **Bounds:** For L-BFGS-B, are `par0` inside bounds, and is the function well-defined throughout the box?
5. **Gradient:** If supplying `gr`, is it correct? If unsure, temporarily drop `gr`.
6. **Method:** Try a different method (e.g. BFGS instead of Nelder–Mead, or vice versa).
7. **Control:** Increase `maxit`, turn on `trace`, and inspect `opt$message`.